

Type-Driven Automated Learning with LALE

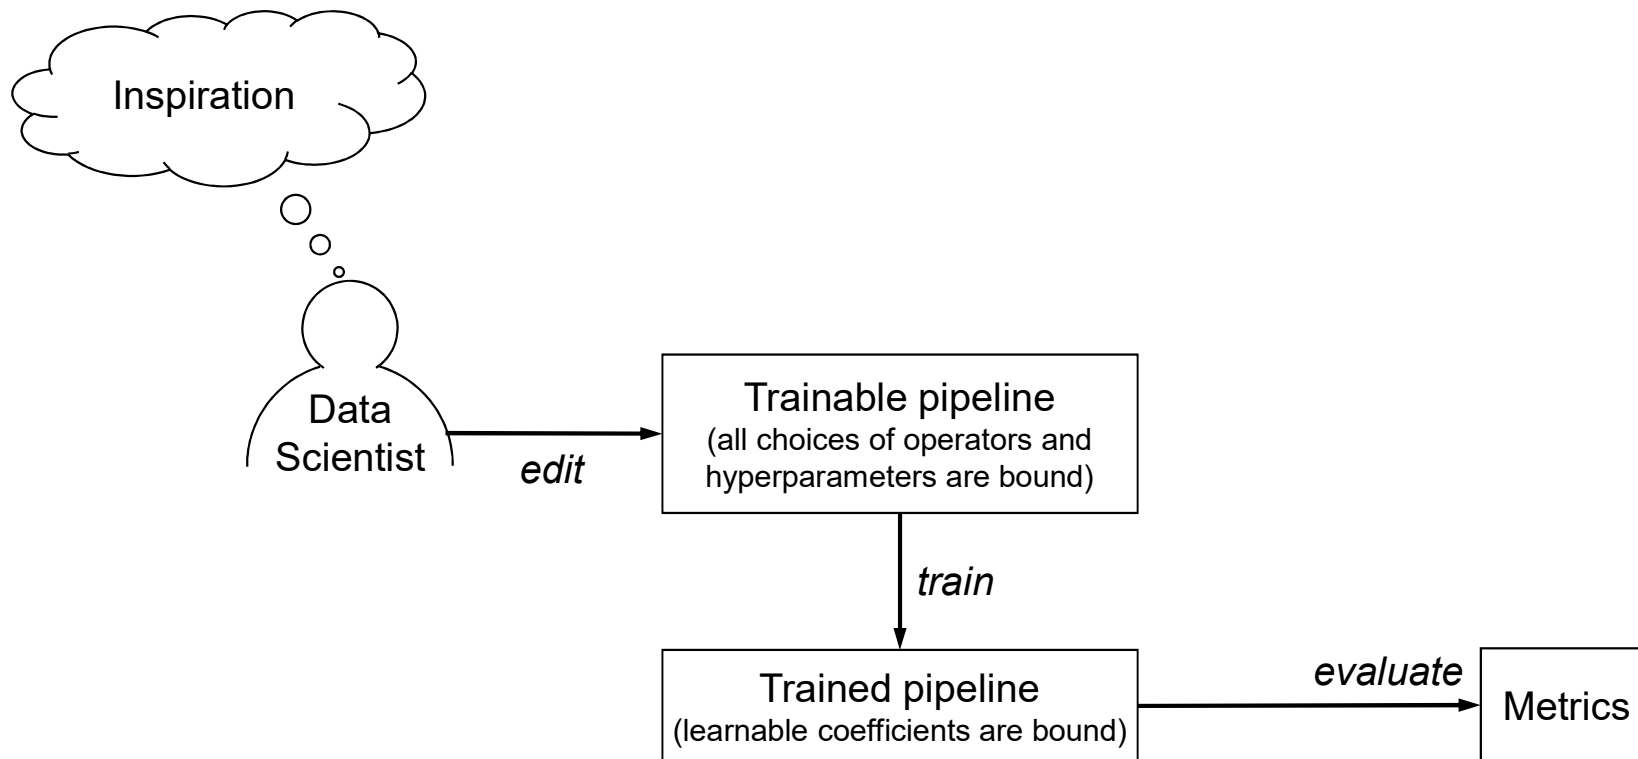
Guillaume Baudart, Martin Hirzel,
Kiran Kate, Pari Ram, and Avi Shinnar

Thursday 9 April 2020

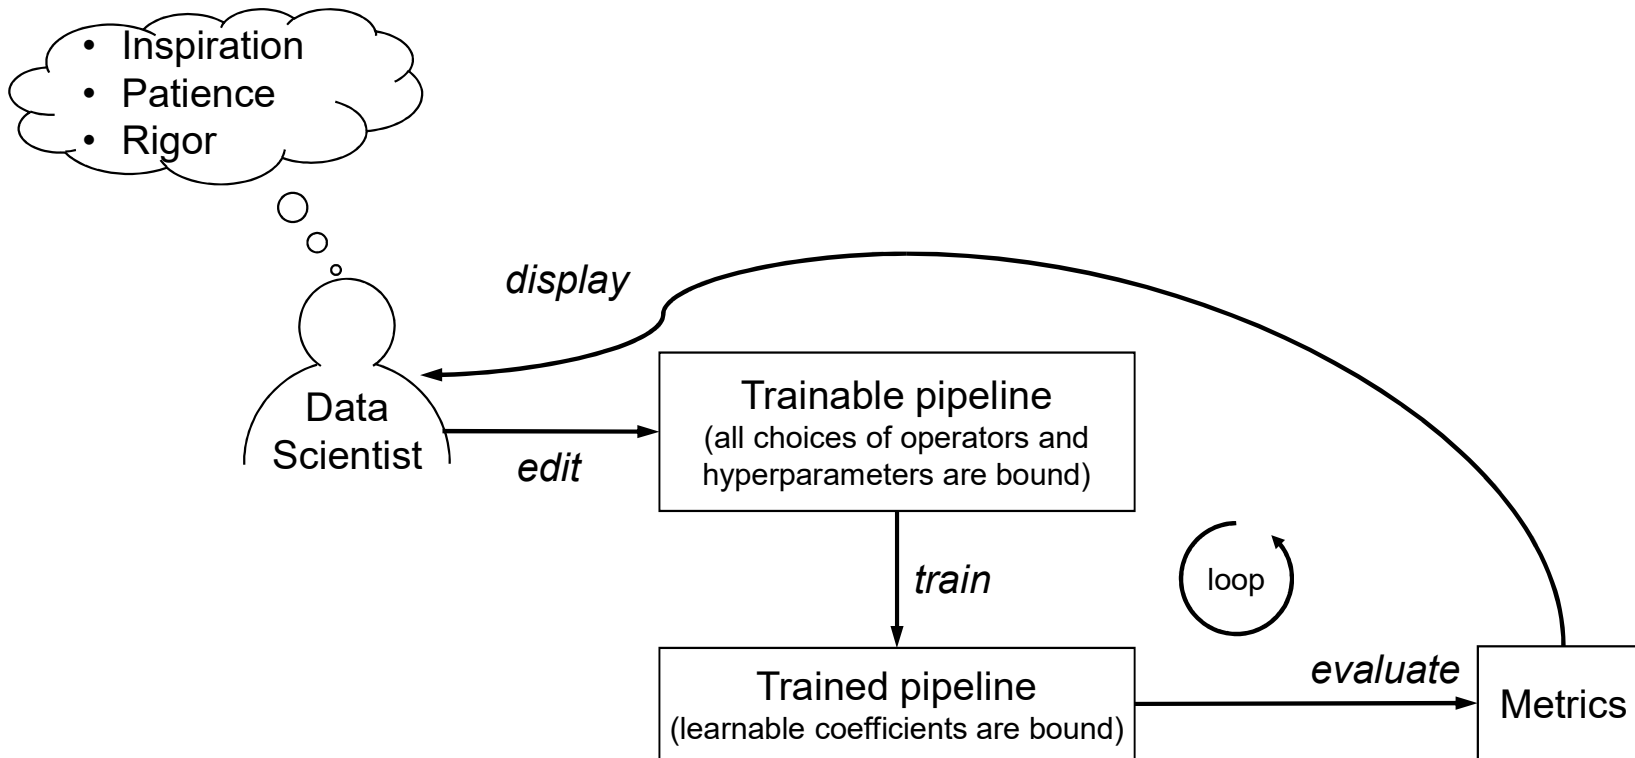
<https://github.com/ibm/lale>



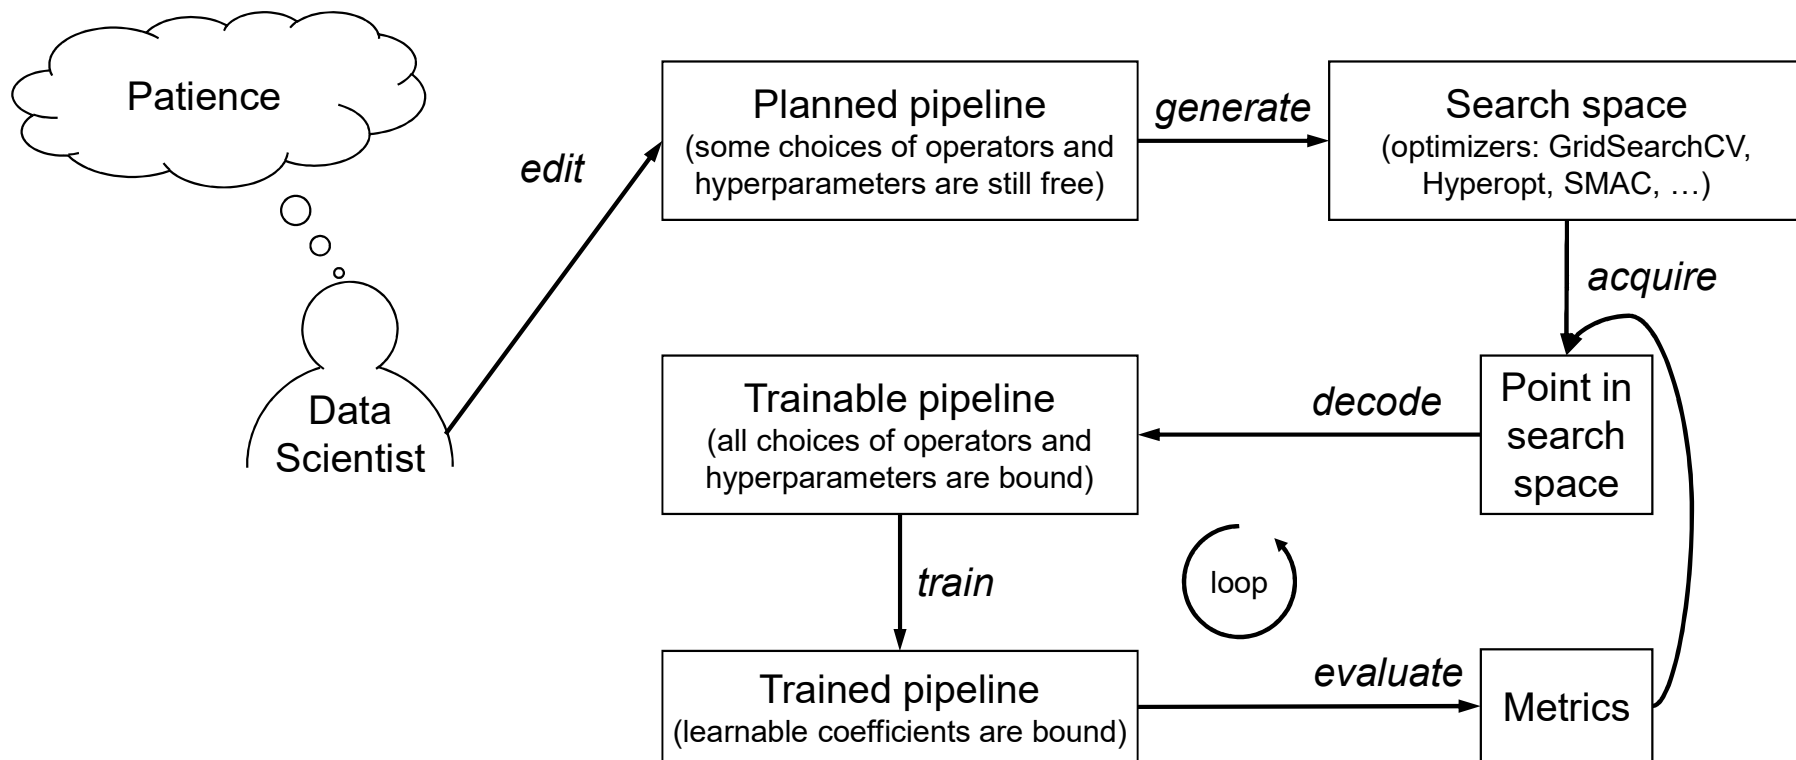
Manual Machine Learning



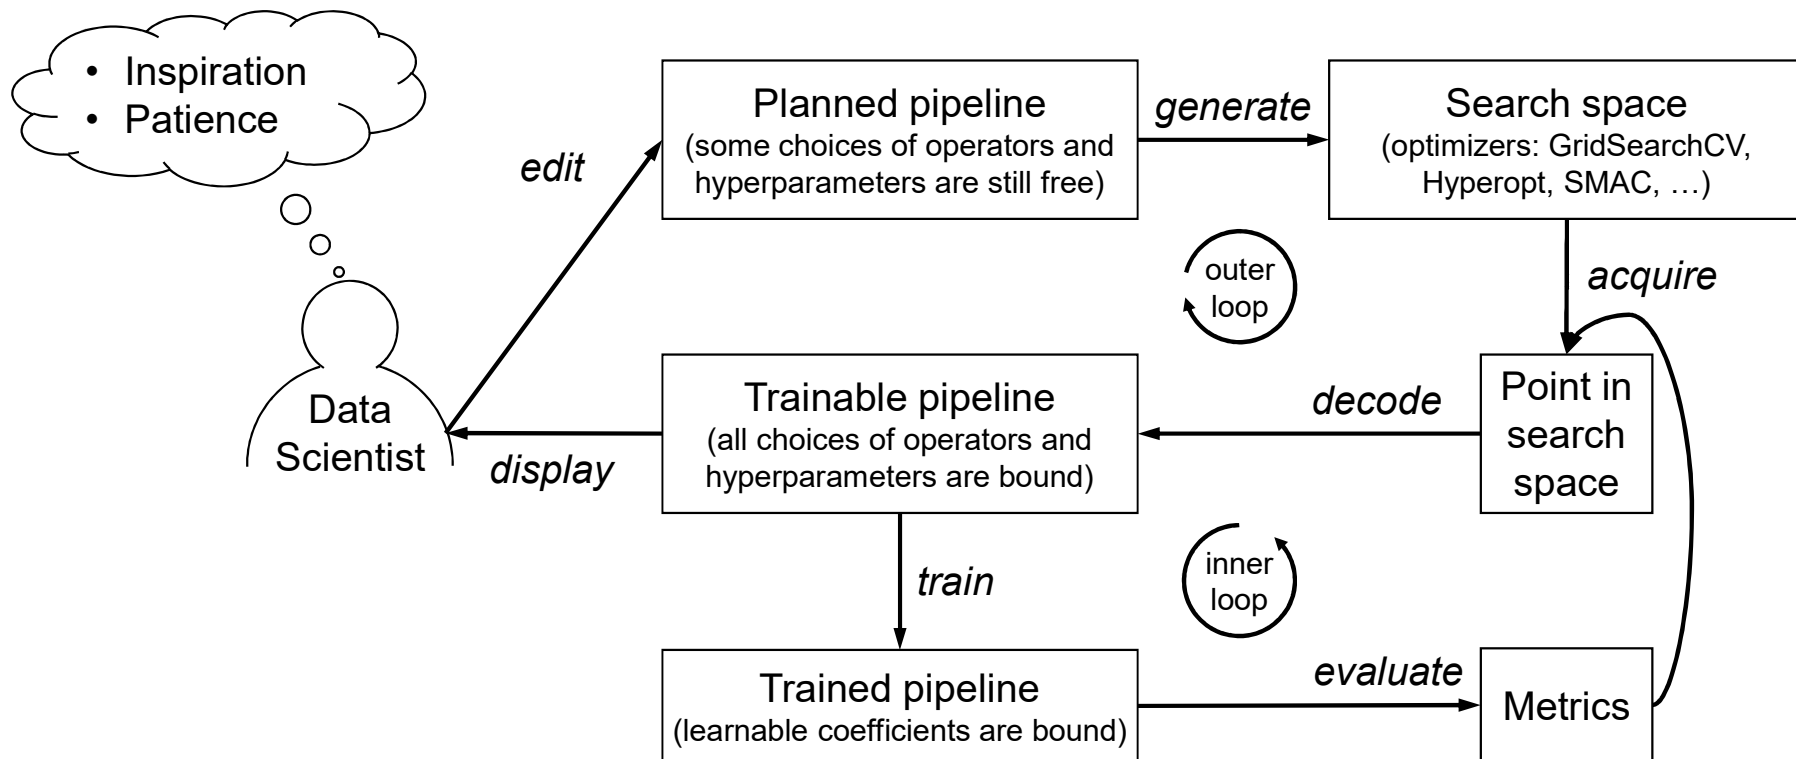
Iterated Manual Machine Learning



AutoML: Automated Machine Learning



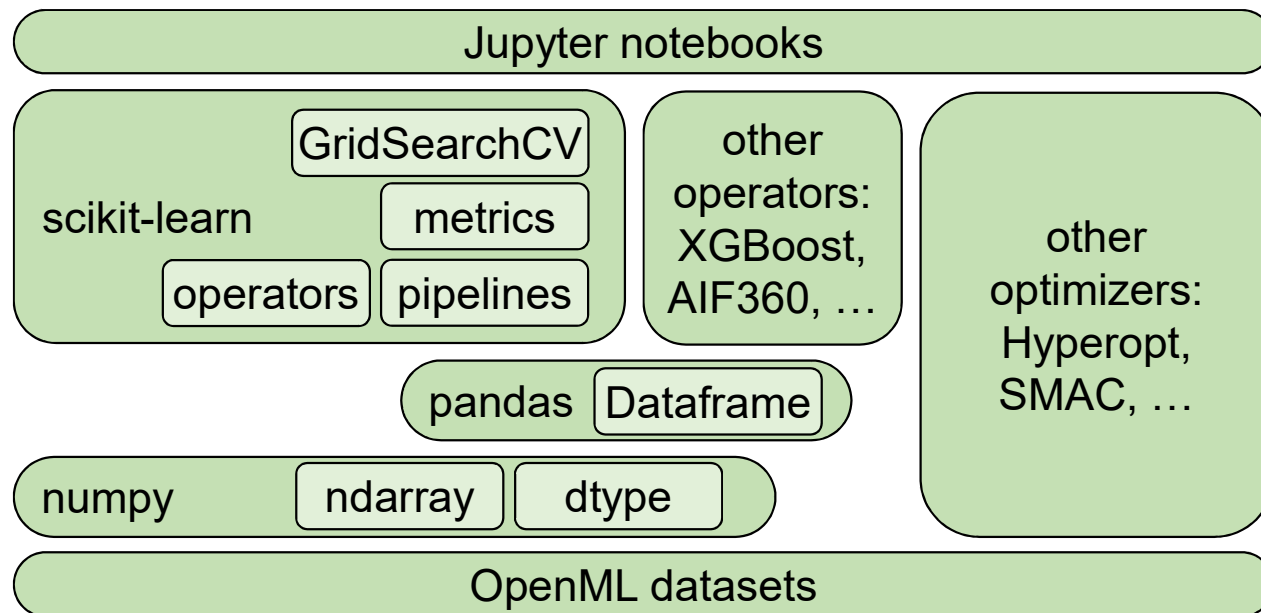
Iterated AutoML



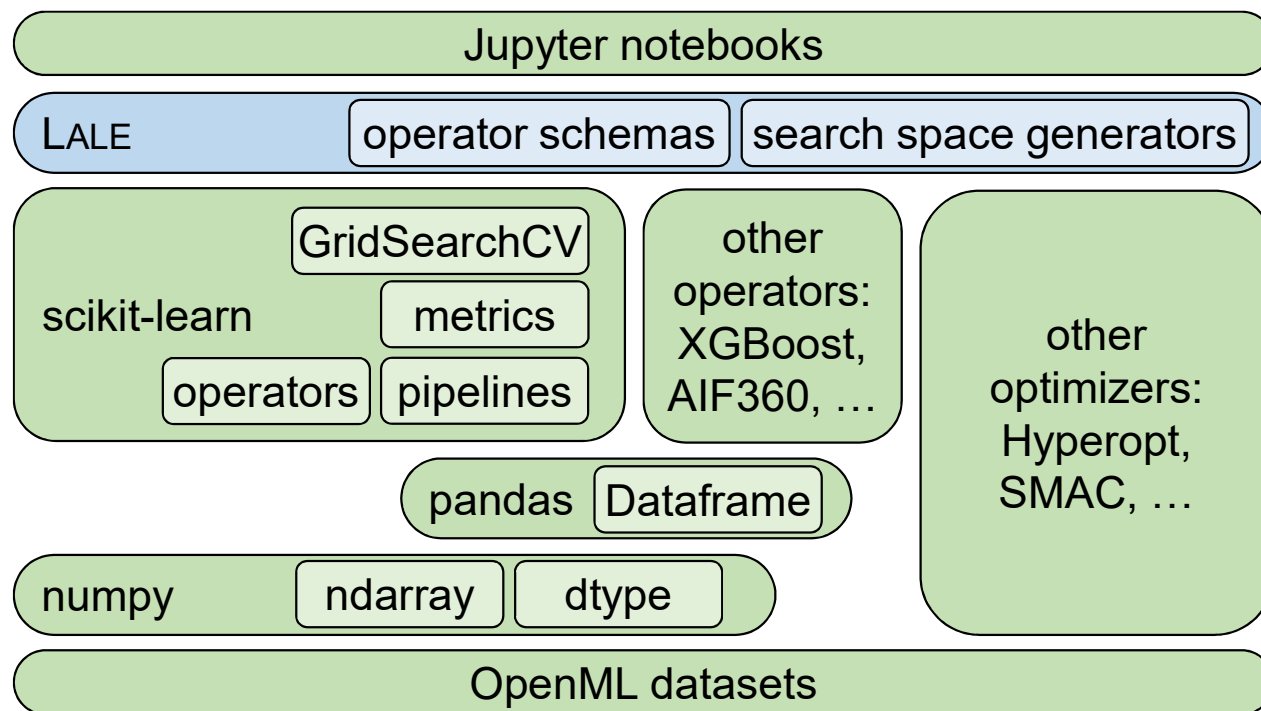
Problem Statement

- Library for semi-automated data science
- Easy-to-use automation
- Consistency of manual vs. automated experience
- Consistency across AutoML optimizers
- Display and iterative refinement
- Expert-level control for human inspiration

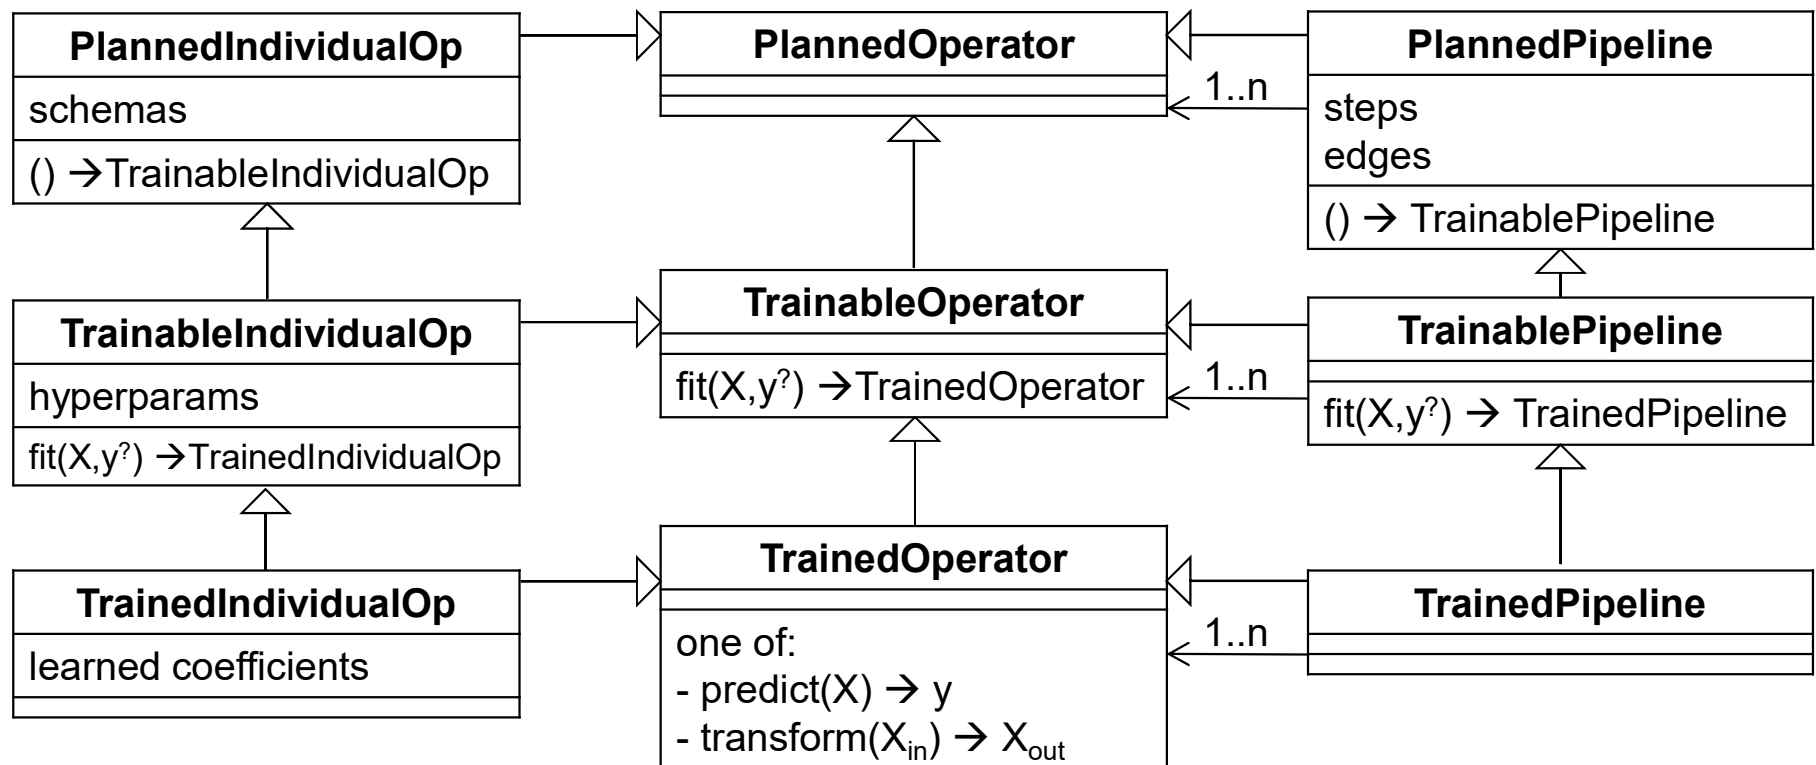
Open-Source AutoML Technologies



LALE's Place in the Open-Source Stack



Operators, Pipelines, and Lifecycle



Legend:
 subtype ("is a") —▷
 attribute ("has a") —>

Example

- Dataset
- Manual machine learning
- Hyperparameter tuning
- Inspecting AutoML results
- Algorithm selection

Example: Covertypes Dataset

```
In [4]: 1 import pandas as pd
2 pd.set_option('display.max_columns', None)
3 pd.concat([pd.DataFrame({'y': train_y}, index=train_X.index),
4            train_X], axis=1).tail(10)
```


Out[4]:

	y	Elevation	Aspect	Slope	Horizontal_Distance_To_Hydrology	Vertical_Distance_To_Hydrology	Horizontal_Distance_To_Roadways	Hillshade_9am	Hillshade_3pm
325384	2	3064.0	86.0	25.0	702.0	259.0	721.0	247.0	253.0
442177	1	3277.0	31.0	15.0	454.0	70.0	1570.0	215.0	245.0
185316	2	3138.0	257.0	14.0	228.0	30.0	5649.0	185.0	236.0
189541	3	2317.0	150.0	8.0	150.0	42.0	644.0	231.0	242.0
428374	2	2970.0	47.0	25.0	319.0	100.0	1919.0	220.0	242.0
234638	1	3278.0	335.0	5.0	360.0	35.0	5763.0	209.0	242.0
172207	1	3175.0	343.0	17.0	162.0	3.0	4395.0	183.0	242.0
240801	1	3355.0	346.0	16.0	180.0	6.0	1922.0	188.0	242.0
435277	1	3154.0	316.0	26.0	339.0	122.0	2688.0	143.0	242.0
297100	7	3344.0	313.0	20.0	0.0	0.0	4317.0	163.0	242.0

Example: Manual Machine Learning

```
In [5]: 1 from sklearn.decomposition import PCA
2 from xgboost import XGBClassifier as XGBoost
3 lale.wrap_imported_operators()
```

```
In [6]: 1 manual_trainable = PCA(n_components=6) >> XGBoost(n_estimators=3)
2 manual_trainable.visualize()
```



```
In [7]: 1 %%time
2 manual_trained = manual_trainable.fit(train_X, train_y)
```

CPU times: user 2.34 s, sys: 1.2 s, total: 3.55 s
Wall time: 2.05 s

```
In [8]: 1 import sklearn.metrics
2 manual_y = manual_trained.predict(test_X)
3 print(f'accuracy {sklearn.metrics.accuracy_score(test_y, manual_y):.1%}')
```

accuracy 64.5%

Example: Hyperparameter Tuning

```
In [9]: 1 XGBoost.hyperparam_schema('n_estimators')
```

```
Out[9]: {'description': 'Number of trees to fit.',
         'type': 'integer',
         'default': 100,
         'minimumForOptimizer': 10,
         'maximumForOptimizer': 1500}
```

```
In [10]: 1 print(PCA.documentation_url())


https://lale.readthedocs.io/en/latest/modules/lale.lib.sklearn.pca.html
```

```
In [11]: 1 from lale.lib.lale import Hyperopt
         2 import lale.schemas as schemas
         3
         4 CustomPCA = PCA.customize_schema(n_components=schemas.Int(min=2, max=54))
         5 CustomXGBoost = XGBoost.customize_schema(n_estimators=schemas.Int(min=1, max=10))
         6
         7 hpo_planned = CustomPCA >> CustomXGBoost
         8 hpo_trainable = Hyperopt(estimator=hpo_planned, max_evals=10, cv=3)
```

```
In [12]: 1 %%time
         2 hpo_trained = hpo_trainable.fit(train_X, train_y)
```

```
100%|██████████| 10/10 [01:20<00:00, 6.64s/trial, best loss: -0.7885106540569516]
CPU times: user 1min 50s, sys: 22.2 s, total: 2min 12s
Wall time: 1min 28s
```

Types as Search Spaces




Example: Inspecting AutoML Results

```
In [13]: ▶ 1 hpo_y = hpo_trained.predict(test_X)
          2 print(f'accuracy {sklearn.metrics.accuracy_score(test_y, hpo_y):.1%}')
```

accuracy 80.1%

```
In [14]: ▶ 1 hpo_trained.get_pipeline().visualize()
```




```
In [15]: ▶ 1 hpo_trained.get_pipeline().pretty_print(ipython_display=True)

from lale.lib.sklearn import PCA
from lale.lib.xgboost.xgb_classifier import XGBoost
import lale
lale.wrap_imported_operators()

pca = PCA(n_components=39, svd_solver='full')
xg_boost = XGBoost(colsample_bylevel=0.6016063807304212, colsample_bytree=0.7763972782064467, learning_rate=0.16389
357351003786, max_depth=10, min_child_weight=5, n_estimators=5, reg_alpha=0.10485915855270356, reg_lambda=0.9268502
695024392, subsample=0.4503841871781402)
pipeline = pca >> xg_boost
```

Example: Algorithm Selection

```
In [19]: ► 1 from sklearn.preprocessing import Normalizer as Norm
2 from sklearn.linear_model import LogisticRegression as LR
3 from sklearn.tree import DecisionTreeClassifier as Tree
4 from sklearn.neighbors import KNeighborsClassifier as KNN
5 from lale.lib.lale import NoOp
6 lale.wrap_imported_operators()
7
8 KNN = KNN.customize_schema(n_neighbors=schemas.Int(min=1, max=10))
9 transp_planned = (Norm | NoOp) >> (Tree | LR(dual=True) | KNN)
10 transp_planned.visualize()
```




Example: Combined Algorithm Selection and Hyperparameter Tuning

```
In [20]: 1 %%time
          2 transp_trained = transp_planned.auto_configure(
          3     train_X, train_y, optimizer=Hyperopt, cv=3, max_evals=3)

100%|██████████| 3/3 [01:48<00:00, 32.59s/trial, best loss: -0.8376392446578157]
CPU times: user 1min 50s, sys: 1.12 s, total: 1min 51s
Wall time: 1min 49s
```

```
In [21]: 1 transp_trained.pretty_print(ipython_display=True, show_imports=False)
          2 transp_trained.visualize()

knn = KNN(algorithm='ball_tree', metric='manhattan', n_neighbors=9)
pipeline = NoOp() >> knn
```



```
In [22]: 1 %%time
          2 transp_y = transp_trained.predict(test_X)
          3 print(f'accuracy {sklearn.metrics.accuracy_score(test_y, transp_y):.1%}')

accuracy 86.6%
CPU times: user 50.6 s, sys: 15.6 ms, total: 50.6 s
Wall time: 50.7 s
```

Expert-Level Control

- Custom metrics
- Non-linear pipelines
- Higher-order operators
- Adding new operators

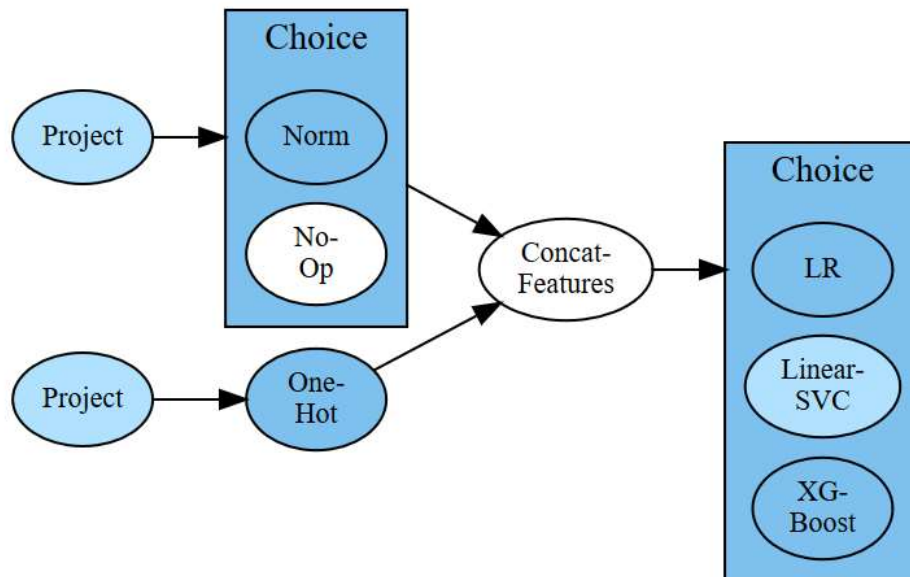
Custom Metrics

```
1 def fairness_scorer(pipeline, X, y):
2     from sklearn.metrics import accuracy_score
3     from aif360.datasets import BinaryLabelDataset
4     predictions = pipeline.predict(X)
5     accuracy = accuracy_score(y, predictions)
6     df = pd.concat([X, y], axis=1)
7     dataset_pred = BinaryLabelDataset(favorable_label=1., unfavorable_label=2.,
8                                       protected_attribute_names=['sex', 'age'],
9                                       df=df, label_names=['credit'])
10    fairness_metrics = aif360.metrics.BinaryLabelDatasetMetric(
11        dataset_pred, unpr_groups, priv_groups)
12    disparate_impact = fairness_metrics.disparate_impact()
13    #Hyperopt minimizes (best_score - score_returned_by_scorer), choosing the values below based on that.
14    if disparate_impact < 0.9 or 1.1 < disparate_impact:
15        return -99
16    else:
17        return accuracy
```

```
1 trained_fairer = planned_fairer.auto_configure(
2     train_X, train_y, optimizer = Hyperopt, cv=3, max_evals=25, scoring=fairness_scorer, best_score=1.0)
```

Non-Linear Pipelines

```
1 project_nums = Project(columns={'type': 'number'})
2 project_cats = Project(columns={'type': 'string'})
3 planned_pipeline = (
4     (project_nums >> (Norm | NoOp) & project_cats >> OneHot)
5     >> ConcatFeatures
6     >> (LR | LinearSVC(dual=False) | XGBoost))
7 planned_pipeline.visualize()
```

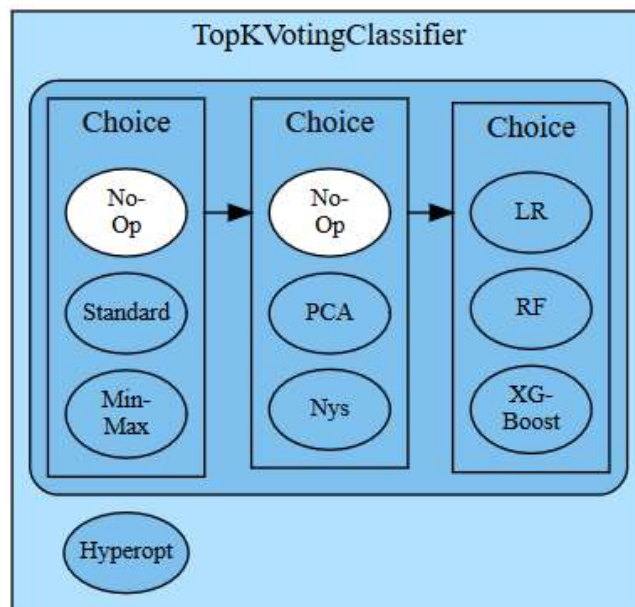


Pipeline Combinators

LAL features	Name	Description	Scikit-learn features
<code>>></code> or <code>make_pipeline</code>	<code>pipe</code>	feed to next	<code>make_pipeline</code>
<code>&</code> or <code>make_union</code>	<code>and</code>	run both	<code>make_union</code> or <code>ColumnTransformer</code>
<code> </code> or <code>make_choice</code>	<code>or</code>	choose one	N/A (specific to given Auto-ML tool)

Higher-Order Operators (1/2)

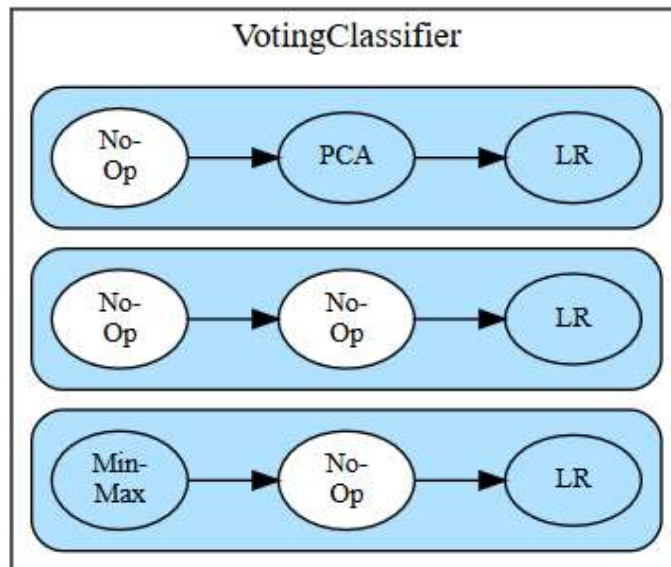
```
In [4]: 1 planned_pipeline = (NoOp | Standard | MinMax) >> (NoOp | PCA | Nys) >> (LR | RF | XGBoost)
2 ensemble = TopKVotingClassifier(
3     estimator=planned_pipeline, k=3, optimizer=Hyperopt,
4     args_to_optimizer={'max_evals':25, 'scoring':'accuracy'})
5 ensemble.visualize()
```



```
In [5]: 1 trained_ensemble = ensemble.fit(train_X, train_y)
```

Higher-Order Operators (2/2)

```
In [8]: 1 best_pipeline = trained_ensemble.get_pipeline()  
2 best_pipeline.visualize()
```



Adding New Operators

- https://nbviewer.jupyter.org/github/IBM/lale/blob/master/examples/docs_new_operators.ipynb
 1. Write class with `__init__`, `fit`, and `predict` or `transform`
 2. Write JSON schemas for hyperparameters and datasets
 3. Call `lale.operators.make_operator`
 4. Write tests for your new operator
- Currently supported operators

Package	Count	Description
<code>lale.lib.sklearn</code>	45	Hand-curated scikit-learn operators
<code>lale.lib.autogen</code>	115	Auto-extracted scikit-learn operators
<code>lale.lib.xgboost</code>	2	Gradient-boosted random forests
<code>lale.lib.aif360</code>	1	Fairness mitigator (incomplete)
... other <code>lale.lib</code>	25	... other packages

Type-Driven Automated Learning

- Types as search spaces
 - E.g., continuous hyperparameter range
 - E.g., categorical hyperparameter enum
- Types as documentation
 - E.g., lale.readthedocs.io/en/latest/modules/lale.lib.sklearn.pca.html
- Types for feature engineering
 - E.g., `Project(columns={'type': 'string'})`
- Types for error checking
 - E.g., datasets, hyperparameters, constraints

Conclusion

- Library for semi-automated data science
 - <https://github.com/ibm/lale>
- Suggested actions:
 - Try it out and send us feedback
 - Contribute new operators
 - Do original research (AutoML optimizers, AI fairness, fast AutoML, ...)